

CSC3310 Algorithms

B-Trees


Matthew, Liam, Helina

Milwaukee School of Engineering

B-Trees

B-Tree

- Long-term (disk, ssd) storage is stored in blocks of data (non-contiguous)
- Most data structures are optimizeddesigned for random-access high-speed ram
- How do we store huge amounts of data efficiently?
- B-trees!
 - Auto-balancing
 - Widened levels&nodes
 - Trees mean pointers




B-Trees

B-Tree Rules

- Nodes cannot have more than a set **Max Degree** number of Keys
- Nodes cannot have fewer than $(\text{Max Degree} // 2)$ Keys
- Child Nodes contain keys between the respective keys of the parent Node

Max Degree: 2



B-Tree Operations

Insert Insert a new element

Delete Remove an element


Search Find an element in the B-Tree

Search

Find an element in the B-Tree

- Scanning each key
 - Recursing through children till key is found.
-
- Algorithm
 - Iterate through subtrees
 - If a key is found, return True
 - If a key is not found, and there is enough children, we recurse through the children.

Max Degree = 4




Insert

Insert a new element

- The data is also the separators (keys)
 - Simply: add to correct spot, split if we violate max degree
- Algorithm
- Step into correct leaf node based off key separators
 - Add data to leaf
 - If we violate max degree
 - Split the values, add the middle value to the parent node
 - Do this step recursively and reassign root if necessary

Max Degree = 4

Insert: 79



Insert

Insert a new element


- The data is also the separators (keys)
- Simply: add to correct spot, split if we violate max degree

Algorithm

- Step into correct leaf node based off key separators
- Add data to leaf
- If we violate max degree
 - Split the values, add the middle value to the parent node
 - Do this step recursively and reassign root if necessary

Max Degree = 4

Insert: 61




Insert

Insert a new element

- The data is also the separators (keys)
 - Simply: add to correct spot, split if we violate max degree
- Algorithm
- Step into correct leaf node based off key separators
 - Add data to leaf
 - If we violate max degree
 - Split the values, add the middle value to the parent node
 - Do this step recursively and reassign root if necessary

Max Degree = 4



Let's look at another example

Insert

Insert a new element


- The data is also the separators (keys)
- Simply: add to correct spot, split if we violate max degree

Algorithm

- Step into correct leaf node based off key separators
- Add data to leaf
- If we violate max degree
 - Split the values, add the middle value to the parent node
 - Do this step recursively and reassign root if necessary

Max Degree = 4

Insert: 2



Insert

Insert a new element


- The data is also the separators (keys)
- Simply: add to correct spot, split if we violate max degree

Algorithm

- Step into correct leaf node based off key separators
- Add data to leaf
- If we violate max degree
 - Split the values, add the middle value to the parent node
 - Do this step recursively and reassign root if necessary

Max Degree = 4

Insert: 3



Insert


Insert a new element

- The data is also the separators (keys)
- Simply: add to correct spot, split if we violate max degree

Algorithm

- Step into correct leaf node based off key separators
- Add data to leaf
- If we violate max degree
 - Split the values, add the middle value to the parent node
 - Do this step recursively and reassign root if necessary

Max Degree = 4



Delete

Remove an element from the tree.

- Need to ensure that the rules are maintained throughout the tree.

Multiple cases

- Case I – Element to be deleted is in a leaf node
- Case II – Element to be deleted is in an internal node
- Case III – The height of the tree shrinks

Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion doesn't violate any properties.
- Deletion results in a violation.


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion doesn't violate any properties.
 - Traverse down the tree to the node
 - Delete the key

Max Degree = 2



Delete 31


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion doesn't violate any properties.
 - Traverse down the tree to the node
 - Delete the key

Max Degree = 2



Delete 31


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion doesn't violate any properties.
 - Traverse down the tree to the node
 - Delete the key

Max Degree = 2



Delete 31


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion doesn't violate any properties.
 - Traverse down the tree to the node
 - Delete the key

Max Degree = 2



Delete 31


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 2



Delete 25


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 2




Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 2



Delete 25


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 2



Delete 25

Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 4



Delete 60

Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 4



Delete 60


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 4



Delete 60


Case I

The element to be deleted is in a leaf node.

Two sub cases:

- Deletion results in a violation.
 - If either sibling has enough keys
 - Borrow from sibling
 - Else
 - Parent node will handle merging

Max Degree = 4



Delete 60

Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child


Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 2



Delete 28


Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 2



Delete 28


Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 2



Delete 28


Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 2




Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 3




Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 3




Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 3




Case II

Element to be deleted is in an internal node

Two sub cases:

- Either child has more than the minimum number of keys
 - Take largest (left) or smallest (right) element from child
- Neither child has more than the minimum number of keys
 - Merge left and right child

Max Degree = 3



Delete 10

Case III

The height of the tree shrinks


- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2




Case III

The height of the tree shrinks

- Deletion of the key leads to fewer number of keys in the node
- If siblings don't have enough keys, we need to merge recursively

Max Degree = 2



Time Complexity

Operation	Time Complexity $n = \# \text{ of nodes in tree}$ $k = \text{max degree}$
Search	$O(\log_k(n) \log(k))$
Insert	$O(\log_k(n) \log(k))$
Delete	$O(\log_k(n) \log(k))$

These can be simplified to $O(\log(n))$ in all cases

$$\log_k(n) = \frac{\log_a(x)}{\log_b(k)}; \log_k(n) * \log(k) = \log(n)$$

Advantages and Disadvantages

Advantages

- Very efficient for huge data storage
 - We set the max degree based on disk block sizes decreasing disk I/O – optimized for disk storage
 - If we are storing huge amounts of data in memory, the larger contiguous nodes result in fewer cache misses
- Auto-Balancing without sacrificing runtime
 - $\log(n)$ for all operations

Disadvantages

- Complex to implement and maintain
 - Memory/Disk Fragmentation
 - Non-numeric data
- A decent amount of overhead no matter the size of the tree.
- Data in range queries can be scattered across nodes
 - B+ Trees are often used which makes this more efficient

Questions?

“Bee Tree”

